

# CS492D: Diffusion Models and Their Applications

## Assignment 3 Session

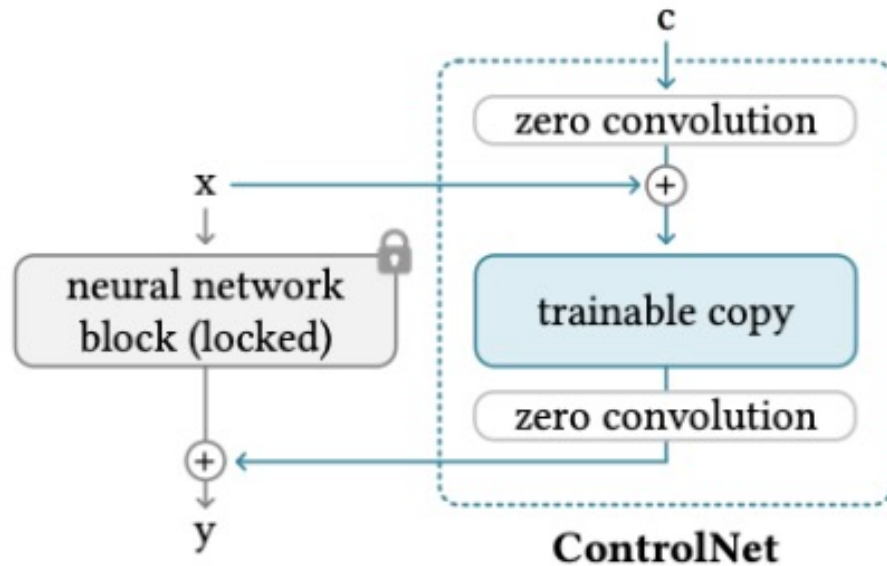
YUSEUNG LEE

Fall 2024  
KAIST


# Introduction

Assignment 3 consists of **two** tasks: ControlNet and LoRA.

## Task 1. ControlNet

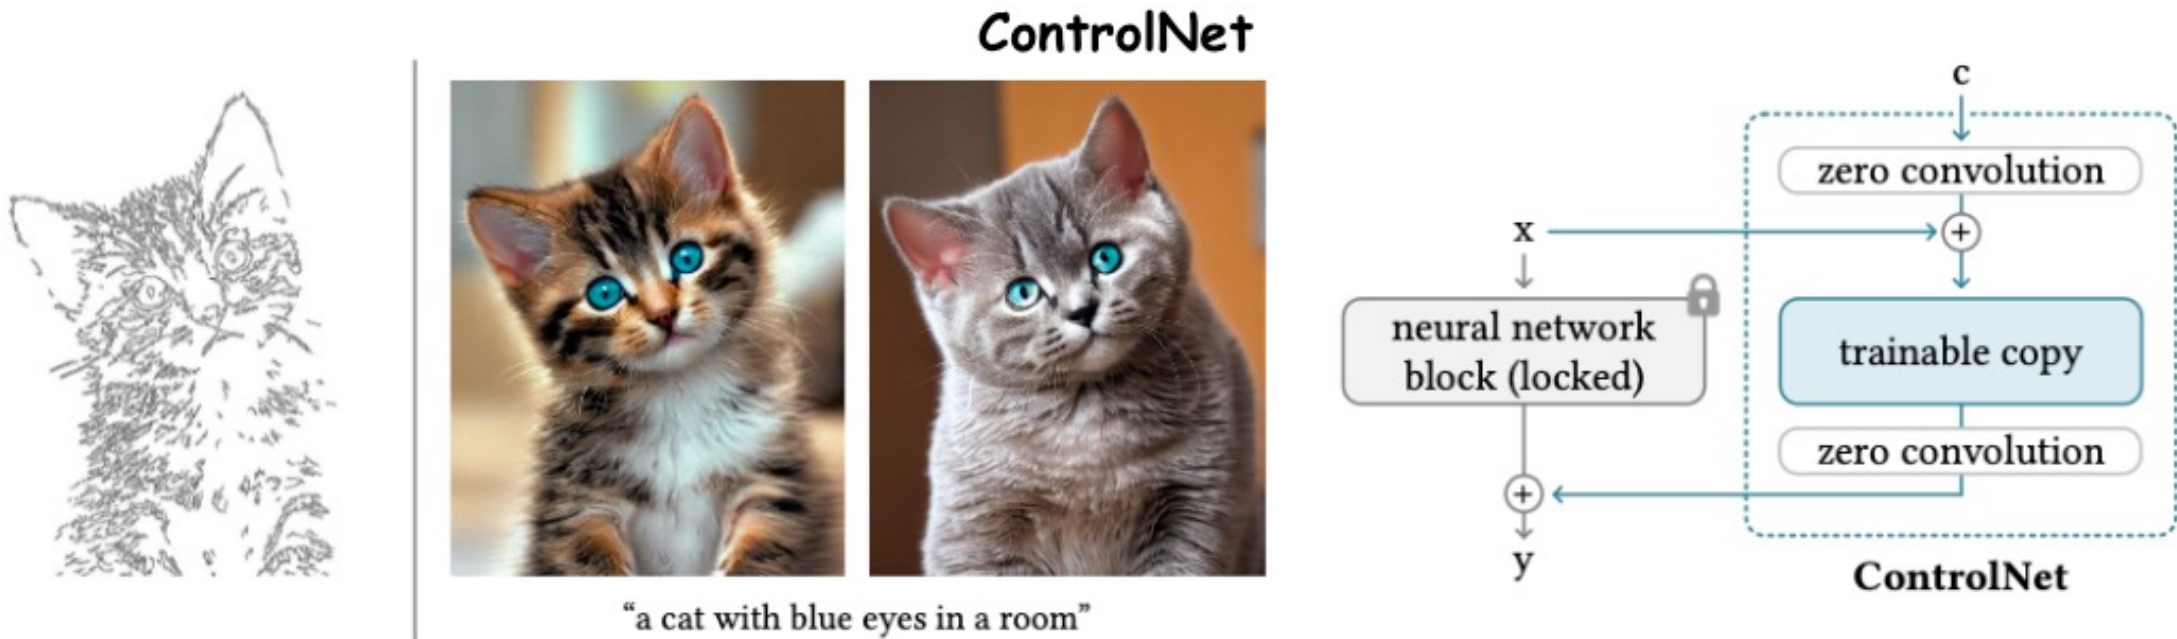


## Task 2. LoRA



# Introduction

In Task 1, you will implement the key components of **ControlNet**.



Adding Conditional Control to Text-to-Image Diffusion Models, Zhang *et al.*, ICCV 2023

# Introduction

In Task 2, you will train custom **LoRA** models with a given library (No implementation).

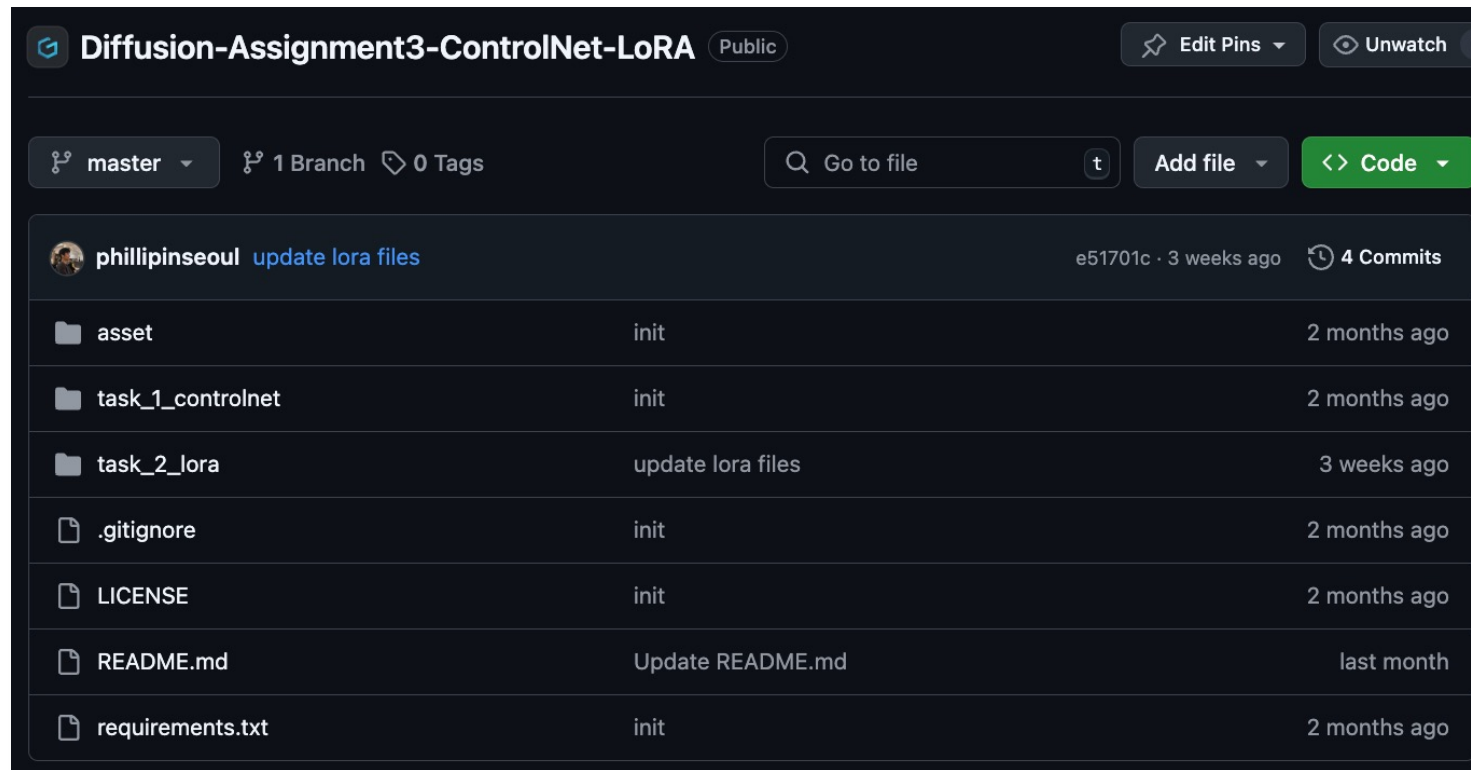


LoRA: Low-Rank Adaptation of Large Language Models, Hu *et al.*, ICLR 2022

# Introduction

The skeleton code and instructions are available at:

<https://github.com/KAIST-Visual-AI-Group/Diffusion-Assignment3-ControlNet-LoRA>



# Important Notes

- All programming assignments are due **two weeks** after the assignment session.
- Late submission will incur **20% penalty** for **each** late day!
- Please carefully check the README of each assignment.
- Missing items in your submission will also incur penalties.

# Overview

0. Introduction to 🤗 Hugging Face and 💣 Diffusers

1. [Task 1] Implementing ControlNet

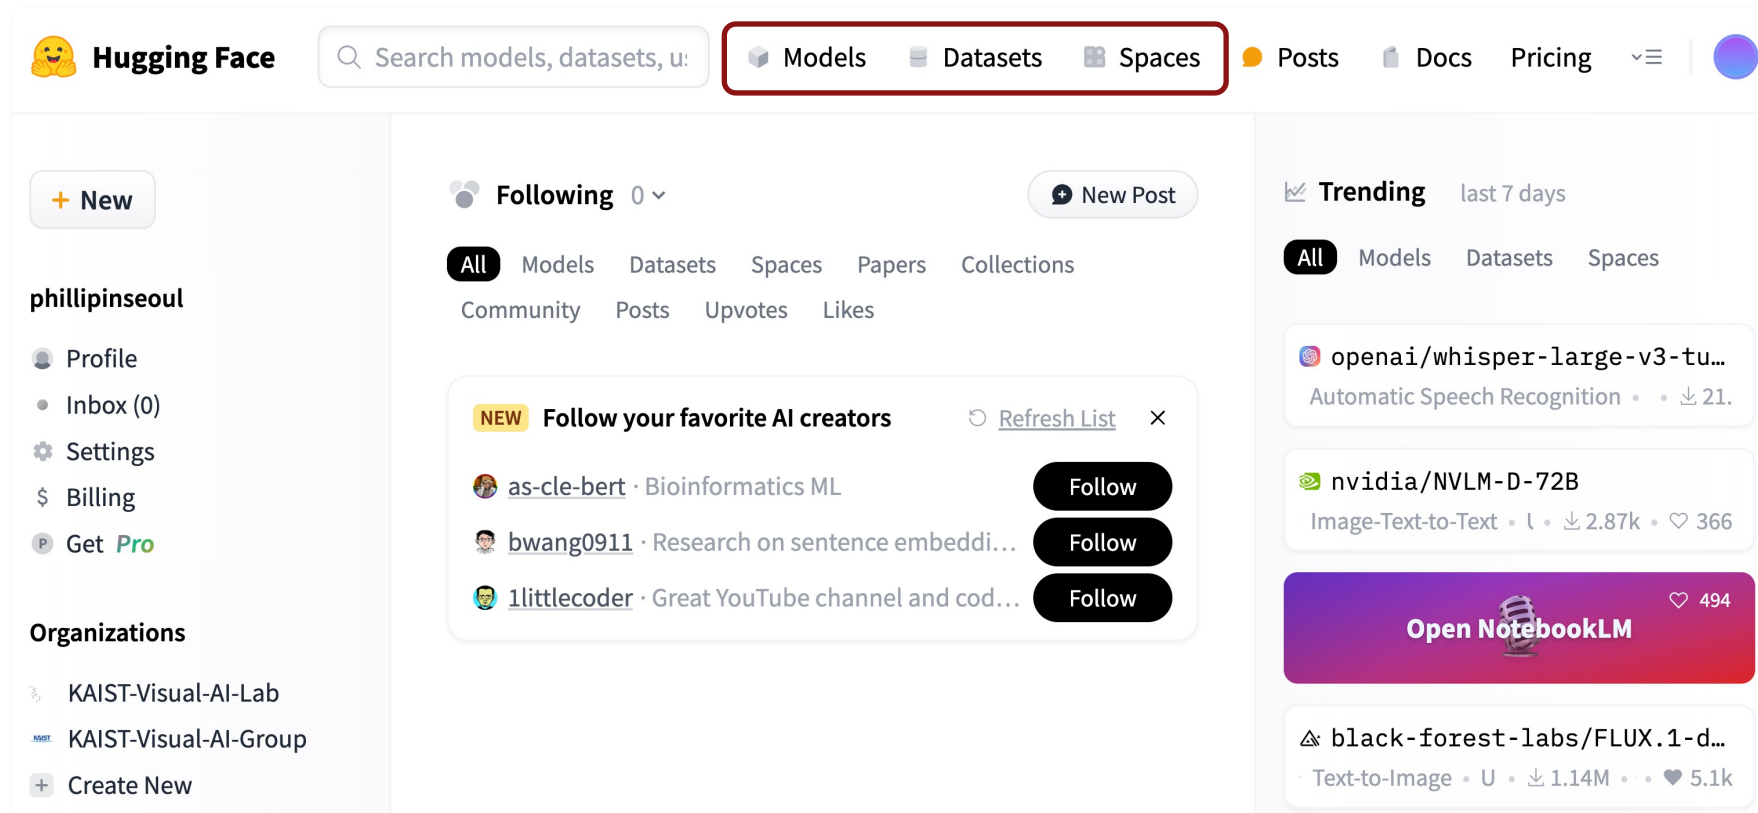
2. [Task 2] Training Custom LoRA Models

# Introduction to Hugging Face and Diffusers



# Hugging Face 🤗

An open-source platform that serves as a hub for machine learning applications.



# Diffusers Library D<sup>🧨</sup>ffusers

Python library for state-of-the-art pre-trained diffusion models.

Popular Tasks & Pipelines		
Task	Pipeline	🤖 Hub
Unconditional Image Generation	DDPM	google/ddpm-ema-church-256
Text-to-Image	Stable Diffusion Text-to-Image	runwayml/stable-diffusion-v1-5
Text-to-Image	unclip	kakaobrain/karlo-v1-alpha
Text-to-Image	DeepFloyd IF	DeepFloyd/IF-I-XL-v1.0
Text-to-Image	Kandinsky	kandinsky-community/kandinsky-2-2-decoder
Text-guided Image-to-Image	Controlnet	lllyasviel/sd-controlnet-canny
Text-guided Image-to-Image	Instruct Pix2Pix	timbrooks/instruct-pix2pix
Text-guided Image-to-Image	Stable Diffusion Image-to-Image	runwayml/stable-diffusion-v1-5
Text-guided Image Inpainting	Stable Diffusion Inpaint	runwayml/stable-diffusion-inpainting
Image Variation	Stable Diffusion Image Variation	lambdalabs/sd-image-variations-diffusers
Super Resolution	Stable Diffusion Upscale	stabilityai/stable-diffusion-x4-upscaler
Super Resolution	Stable Diffusion Latent Upscale	stabilityai/sd-x2-latent-upscaler

e.g. Loading Stable Diffusion from diffusers

```
import torch
from diffusers import StableDiffusionPipeline

model_id = "runwayml/stable-diffusion-v1-5"
pipe = StableDiffusionPipeline.from_pretrained(
    model_id,
    torch_dtype=torch.float16
)
pipe = pipe.to("cuda")

prompt = "a photo of an astronaut riding a horse on mars"
image = pipe(prompt).images[0]

image.save("astronaut_rides_horse.png")
```

# What to Do: Task 0

You need access to Hugging Face to proceed with Task 1 and 2:

1. Sign into Hugging Face.
2. Obtain your access token at <https://huggingface.co/settings/tokens>.
3. From your terminal, log into Hugging Face using

```
$ huggingface-cli login
```

and enter your Access Token.

# What to Do: Task 0

4. To check the access to Hugging Face, download Stable Diffusion from Hugging Face and generate an image with it:

```
import torch
from diffusers import StableDiffusionPipeline

model_id = "CompVis/stable-diffusion-v1-4"
device = "cuda"

pipe = StableDiffusionPipeline.from_pretrained(model_id, torch_dtype=torch.float16)
pipe = pipe.to(device)

prompt = "a photo of an astronaut riding a horse on mars"
image = pipe(prompt).images[0]

image.save("astronaut_rides_horse.png")
```


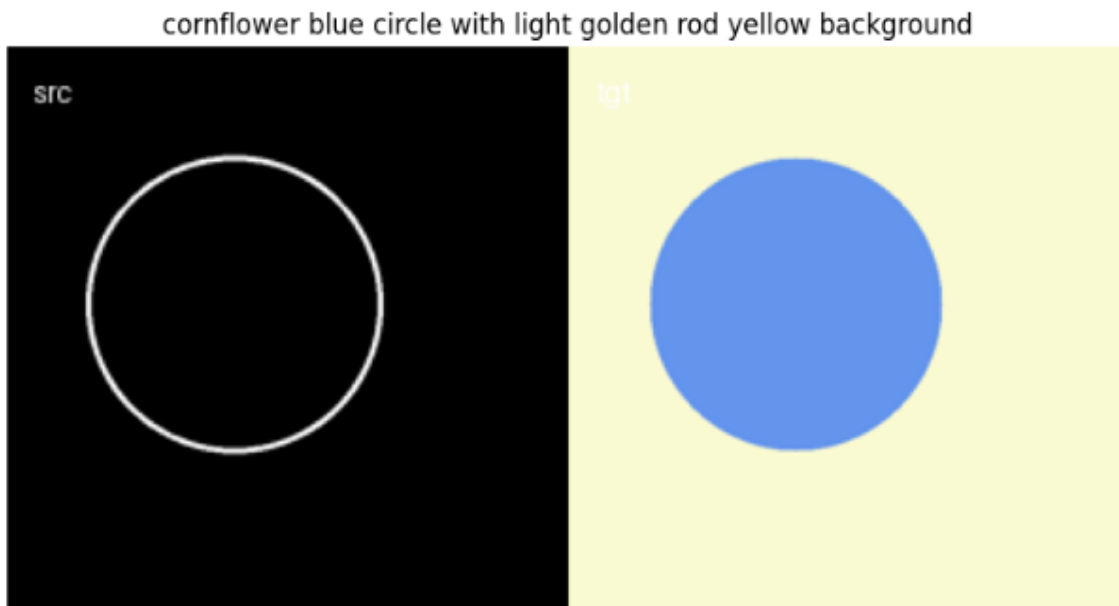
# **[Task 1]**

## **Implementing ControlNet**

# What to Do: Task 1

**Goal:** Train ControlNet on Fill50K dataset.

Fill50K

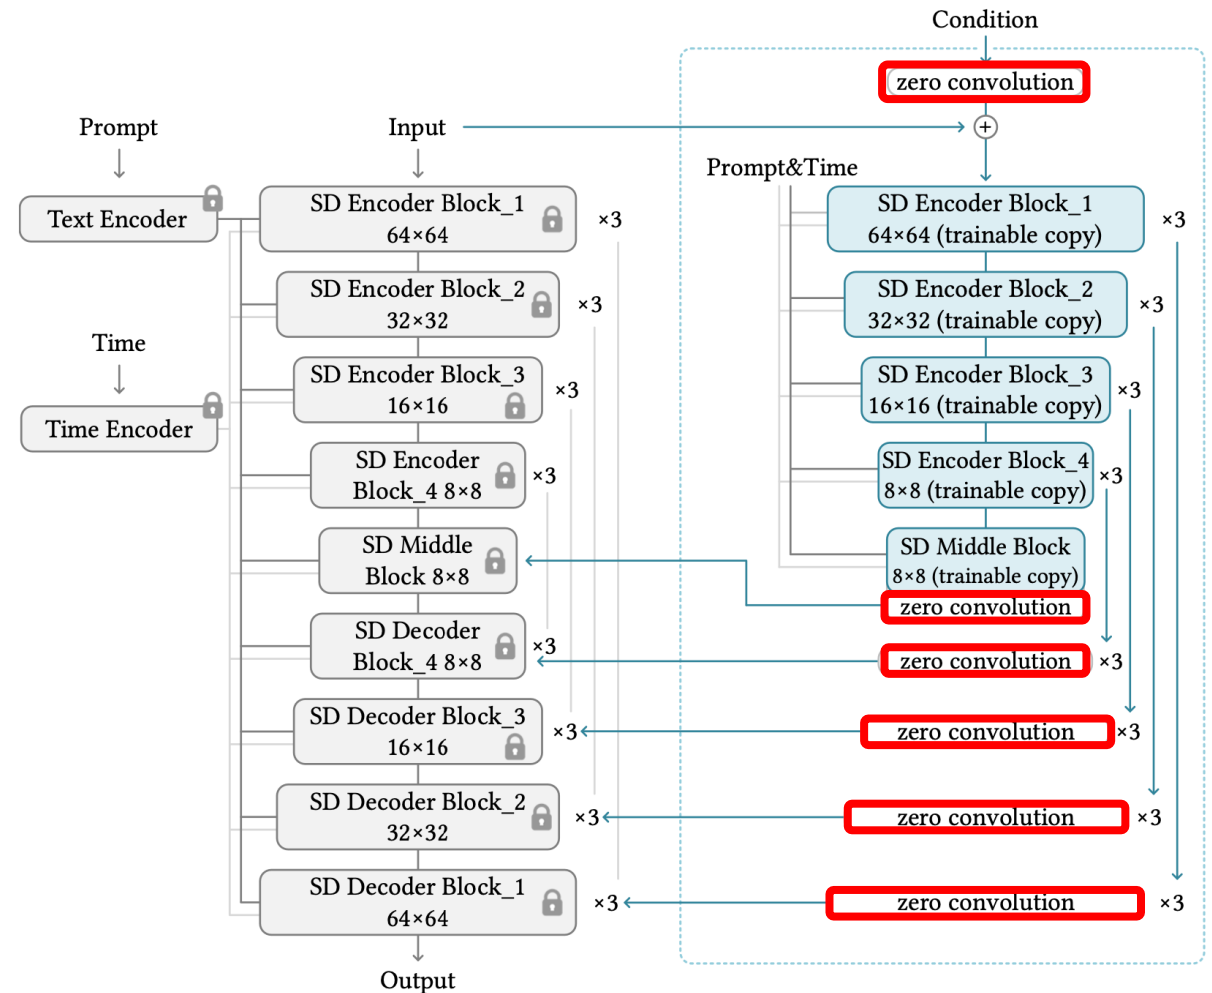


# What to Do: Task 1

## Implement Zero-Convolution.

```
def zero_convolution(
    in_channels, out_channels,
    kernel_size=1, stride=1, padding=0
):
    """
    Return a 'zero-convolution layer'
    (Initialized weight & bias as zeros.)
    """
    ##### TODO (1) #####
    # DO NOT change the code outside this part.
    # Return a zero-convolution layer,
    # with the weight & bias initialized as zeros.
    module = None
    ##### TODO (1) #####
    return module
```

task\_1\_controlnet/diffusion/controlnet.py

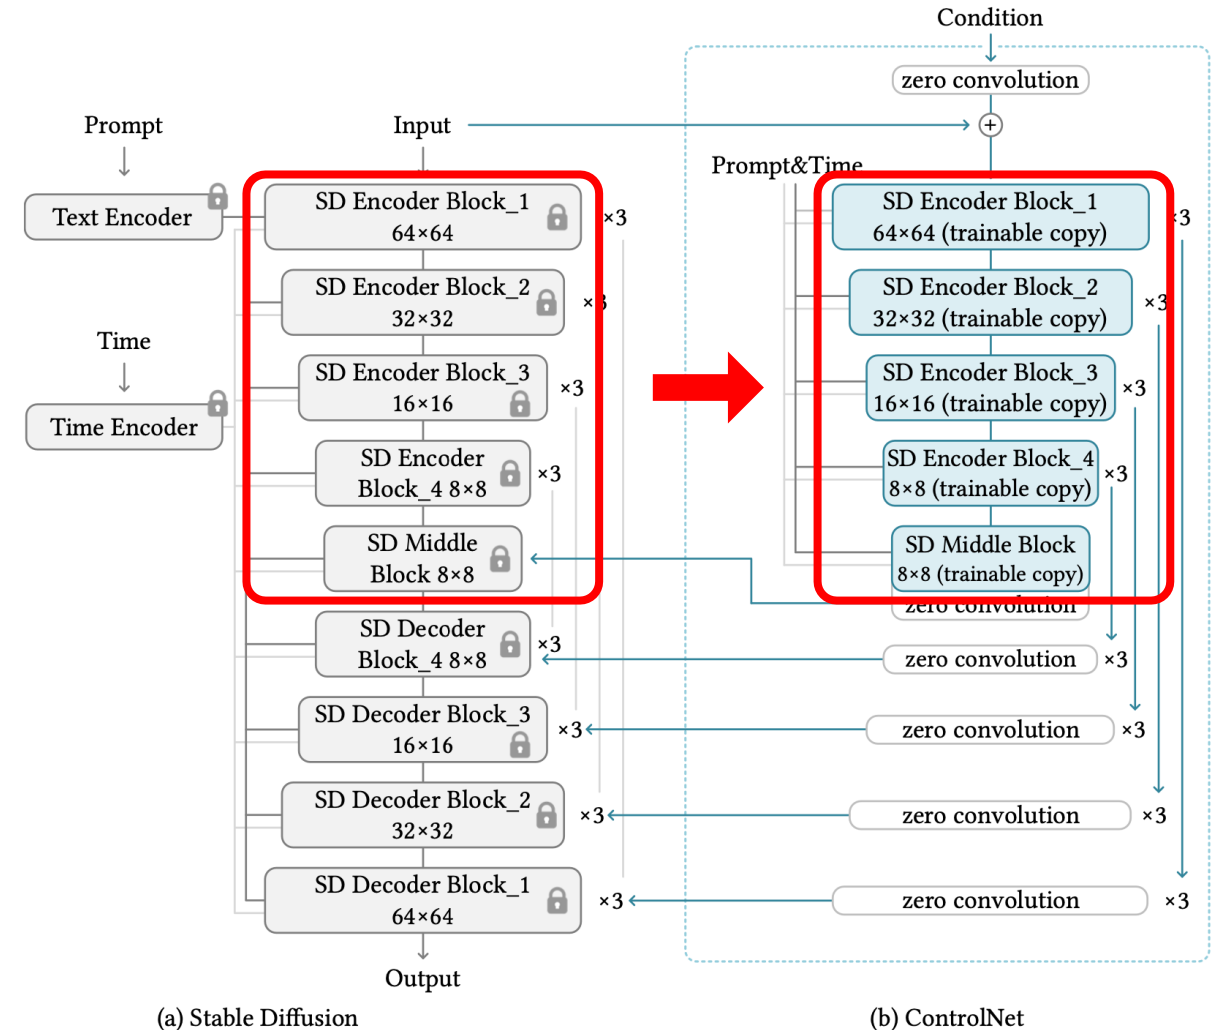


(a) Stable Diffusion

(b) ControlNet

# What to Do: Task 1

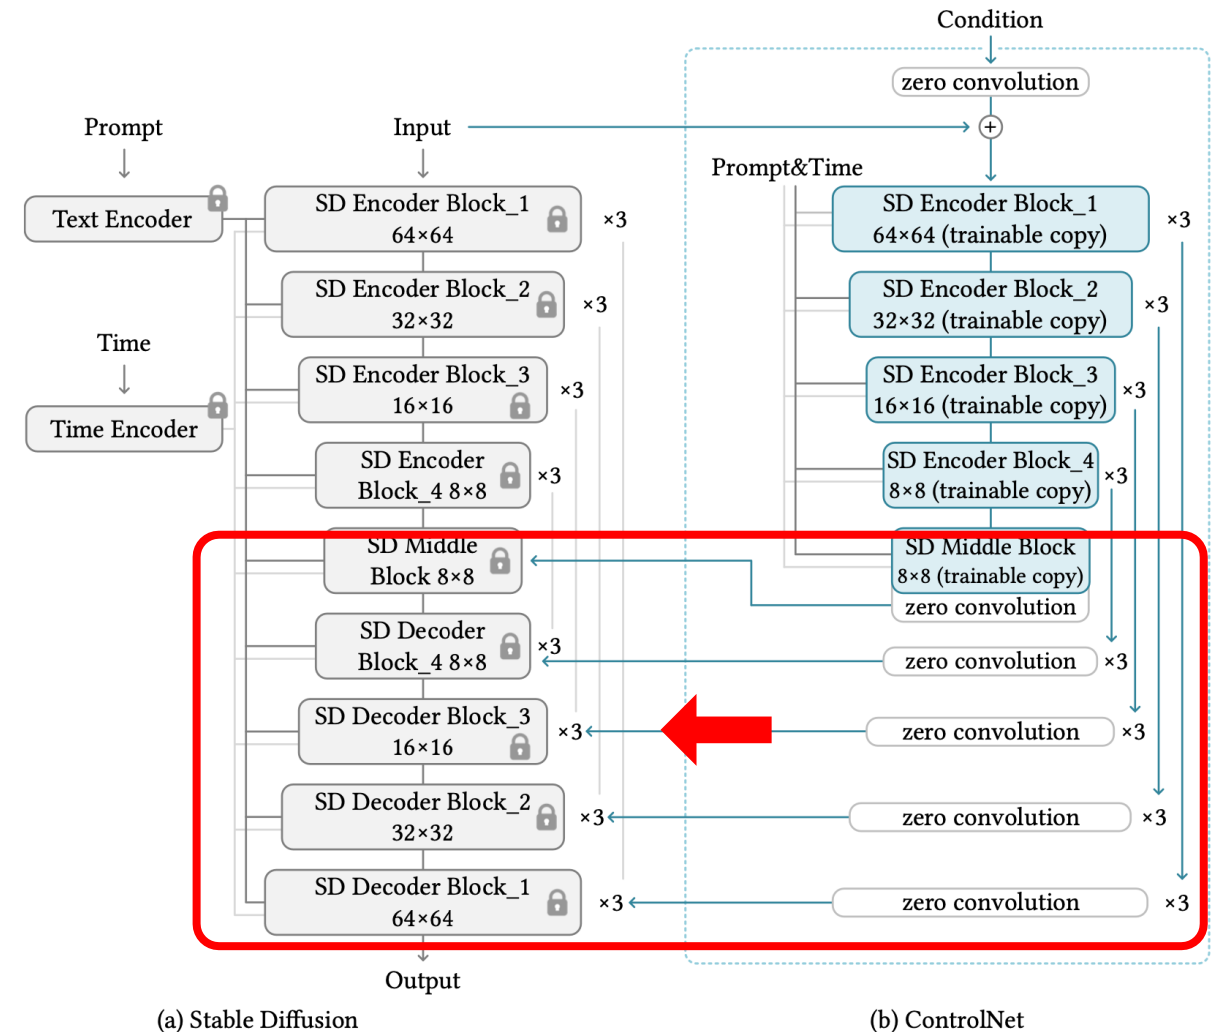
Initialize ControlNet using a pre-trained U-Net from Stable Diffusion.





# What to Do: Task 1


- Apply Zero-Convolution to the residual features of each ControlNet block.
- Integrate outputs from ControlNet blocks into U-Net decoder of Stable Diffusion.



# What to Do: Task 1

- Train ControlNet using: `$ sh train.sh`
- Generate images with 5 different conditions (./data/test\_conditions) and text prompts from (./data/text\_prompts.json).

Test conditions




Test prompts

```
{  
  "0": "pale golden rod circle with old lace background",  
  "1": "sea green circle with a light cyan background",  
  "2": "deep sky blue circle with a light yellow background",  
  "3": "rosy brown circle with a misty rose background",  
  "4": "forest green circle with an antique brown background"  
}
```

# What to Do: Task 1 (Optional)

Train ControlNet on a different type of condition (e.g. depth map, canny edge, sketch, etc.).



You can use either an **open-source dataset** or your own **custom data**.



Canny Edge



Depth (midas)



# **[Task 2]**

## **Training Custom LoRA Models**



# What to Do: Task 2

Goal: Train custom LoRA models on three different datasets.

Training Data



Generated with GPT4

Before LoRA



After LoRA



*“A man with sunglasses”*

# What to Do: Task 2

**No implementation required!**

The main objective of this task is to:

- **Gain hands-on experience** on customizing diffusion models using LoRA,
- **Explore creative use-cases** of state-of-the-art diffusion models.

# What to Do: Task 2

[Task 2-1] Train LoRA on a specific **style**.

You can either use an **open-source dataset** and train with

```
$ sh scripts/train_lora.sh
```

Or create a **custom dataset** and train with

```
$ sh scripts/train_lora_custom.sh
```

# What to Do: Task 2

[Task 2-1] Train LoRA on a specific **style**.

In either case, just simply set the path to the dataset by

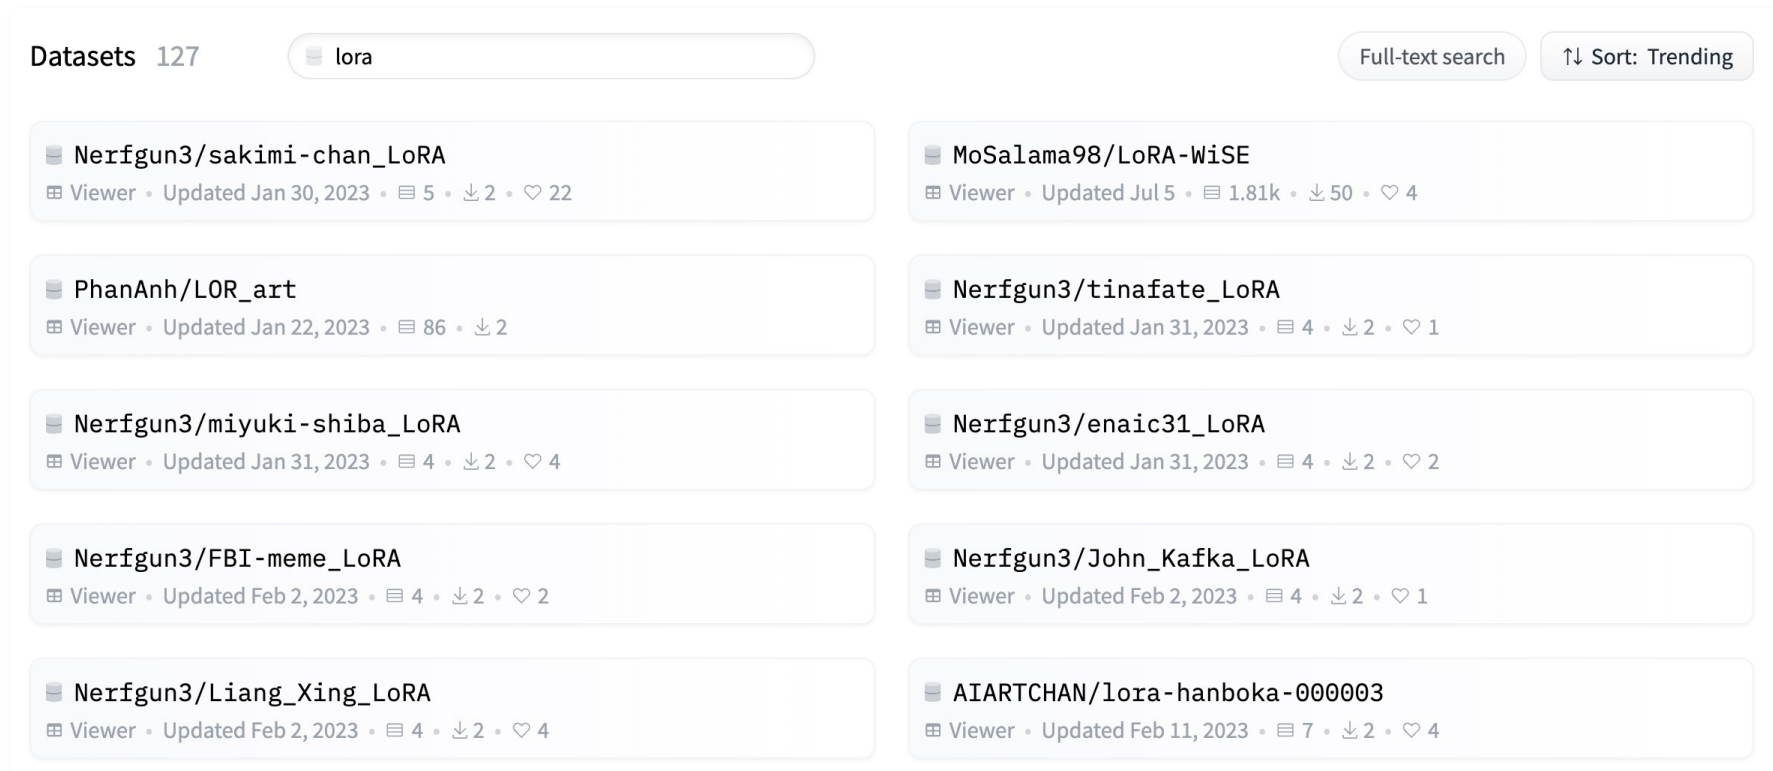
```
$ export DATASET_NAME="$PATH_TO_DATASET"
```

in the given script file.



# What to Do: Task 2

**TIP:** You can find many open-source LoRA training datasets on Hugging Face.



[https://huggingface.co/datasets?modality=modality:image&size\\_categories=or:%28size\\_categories:n%3C1K,size\\_categories:1K%3Cn%3C10K%29&sort=trending&search=lora](https://huggingface.co/datasets?modality=modality:image&size_categories=or:%28size_categories:n%3C1K,size_categories:1K%3Cn%3C10K%29&sort=trending&search=lora)

# What to Do: Task 2

[Task 2-2] Train LoRA on a specific **identity** using DreamBooth + LoRA.

Run `$ sh scripts/train_dreambooth_lora.sh` for training.

Training Data



Generated with GPT4

Before LoRA



After LoRA



# What to Do: Task 2

Inference Code: `lora_inference.ipynb`

## Load Stable Diffusion

```
In [ ]: pipe = StableDiffusionPipeline.from_pretrained(
        "CompVis/stable-diffusion-v1-4",
        torch_dtype=torch.float16
    )
print("[INFO] Successfully loaded Stable Diffusion!")
```

## Load LoRA weights

Change to your LoRA weights!

```
In [22]: # lora_path = "./runs/sd-naruto-model-lora"
lora_path = "./runs/artistic_custom"
# lora_path = "./runs/dreambooth_cat"
# lora_path = None # if not using LoRA

if lora_path is not None:
    pipe.load_lora_weights(lora_path)
    print("[INFO] Successfully loaded LoRA weights!")

pipe = pipe.to(device)
```

[INFO] Successfully loaded LoRA weights!

## Inference

Change prompt and seed for diverse outputs!

```
In [27]: prompt = "a man with sunglasses"
seed = 10

seed_everything(seed)

image = pipe(
    prompt,
    num_inference_steps=30,
    guidance_scale=7.5
).images[0]

image
```

100% |██████████| 30/30 [00:01<00:00, 22.53it/s]

# What to Submit

Submit a zip file named {NAME}\_{STUDENT\_ID}.zip that includes:

- Code for **Task 1**
  - Include your code inside `task_1_controlnet/` directory.
- LoRA Checkpoints for **Task 2**
  - `pytorch_lora_weights.safetensors`
- PDF Report (Max. 2 pages)

# What to Submit

The PDF report should include:

## [Task 1]

- 5 different condition inputs, corresponding text prompts, and the generated images.
- A brief analysis of the results for each condition.
- **(Optional Task 1-1)** 5 different condition inputs, corresponding text prompts, and the generated images.
- **(Optional Task 1-1)** A brief explanation about the training dataset and the training results.

# What to Submit

The PDF report should include:

## **[Task 2]**

- (Task 2-1) Description on the dataset used, including its source.
- (Task 2-1) Visualization of training images and generated image with the corresponding text prompts.
- (Task 2-2) Description on the dataset used, including its source.
- (Task 2-2) Visualization of training images and generated image with the corresponding text prompts.

# What to Submit

The zip file should look like:

```
.
├── 2024XXXX.pdf          <-- report (max. 2 pages)
├── task_1_controlnet     <-- code for Task 1
├── lora_1                 <-- checkpoints for Task 2
│   └── pytorch_lora_weights.safetensors
├── lora_2
│   └── pytorch_lora_weights.safetensors
└── lora_3
    └── pytorch_lora_weights.safetensors
```

# Grading

The scores for each task are detailed as follows:

## Task 1 (10pt):

- [0pt] Either the code or the report is not submitted.
- [5pt] Generated images do not align with the input conditions.
- [10pt] Generated images accurately align with the input conditions.



# Grading

The scores for each task are detailed as follows:

## Task 2 (10pt):

- [0pt] The report is not submitted.
- [5pt] Outputs of either one of the LoRAs do not align with the training data.
- [10pt] Outputs of both LoRAs accurately align with the training data.

# ! Disclaimer !

Ethical usage of AI models is important:

1. You can freely refer to any online source code when working on Assignment 3. BUT we prohibit simply copying & pasting existing code. The aim for the assignment is to get hands-on experience on how the code works for diffusion model applications.
2. You **MUST NOT** use personalization techniques for **unethical purposes**, such as generating content that includes nudity, violence of specific identities. It is your responsibility to ensure that these methods are applied ethically.

# ! Disclaimer !

**Ethical usage** of AI models is important:

3. When using open-source datasets, use **MUST** give credit to the original providers. If you are using such datasets for training ControlNet or LoRA, remember to clearly cite the original source of data.

# Demo

# Thank You