

CS492D: Diffusion Models and Their Applications


Assignment 1 Session

SEUNGWOO YOO

Fall 2024
KAIST

Introduction

In Assignment 1, you will implement the key components of **Denoising Diffusion Probabilistic Models (DDPMs)**.



Denoising Diffusion Probabilistic Models, Ho *et al.*, NeurIPS 2020

Introduction

The skeleton code and instructions are available at:

<https://github.com/KAIST-Visual-AI-Group/Diffusion-Assignment1-DDPM>

The screenshot shows the GitHub repository page for 'Diffusion-Assignment1-DDPM'. The repository is public and has 3 stars, 1 watch, and 3 forks. The main branch is 'main'. The repository was created by 'DeveloperY0115' 4 days ago, with 15 commits. The file list includes: '2d_plot_diffusion_todo' (Remove DDIM sampling, 4 days ago), 'assets' (Init repo, 2 weeks ago), 'image_diffusion_todo' (Reduce batch size during sampling, 5 days ago), '.gitignore' (Init repo, 2 weeks ago), 'LICENSE' (Init repo, 2 weeks ago), 'README.md' (Update README.md, last week), and 'requirements.txt' (Init repo, 2 weeks ago). The README section is titled 'Denoising Diffusion Probabilistic Models (DDPM)' and includes the following text: 'KAIST CS492(D): Diffusion Models and Their Applications (Fall 2024) Programming Assignment 1', 'Instructor: Minhyuk Sung (mhsung [at] kaist.ac.kr)', 'TA: Seungwoo Yoo (dreamy1534 [at] kaist.ac.kr)', and 'Credit: Juil Koo (63days [at] kaist.ac.kr) & Nguyen Minh Hieu (hieuristics [at] kaist.ac.kr)'. The right sidebar shows the 'About' section with no description, website, or topics provided, and the 'Releases' section with no releases published.

Diffusion-Assignment1-DDPM Public

Edit Pins Watch 1 Fork 3 Star 3

main 1 Branch 0 Tags

Go to file Add file Code

DeveloperY0115 Remove DDIM sampling 2cf8323 · 4 days ago 15 Commits

2d_plot_diffusion_todo	Remove DDIM sampling	4 days ago
assets	Init repo	2 weeks ago
image_diffusion_todo	Reduce batch size during sampling	5 days ago
.gitignore	Init repo	2 weeks ago
LICENSE	Init repo	2 weeks ago
README.md	Update README.md	last week
requirements.txt	Init repo	2 weeks ago

README License

Denoising Diffusion Probabilistic Models (DDPM)

[KAIST CS492\(D\): Diffusion Models and Their Applications \(Fall 2024\)](#)
Programming Assignment 1

Instructor: [Minhyuk Sung](#) (mhsung [at] kaist.ac.kr)
TA: [Seungwoo Yoo](#) (dreamy1534 [at] kaist.ac.kr)
Credit: [Juil Koo](#) (63days [at] kaist.ac.kr) & [Nguyen Minh Hieu](#) (hieuristics [at] kaist.ac.kr)

About

No description, website, or topics provided.

Readme View license Activity Custom properties 3 stars 1 watching 3 forks Report repository

Releases

No releases published [Create a new release](#)

Packages

No packages published [Publish your first package](#)

Contributors 2

[DeveloperY0115](#) SeungWoo Yoo
[mhsung](#) Minhyuk Sung

Introduction

- All programming assignments are due **two weeks** after the assignment session.
- Late submission will incur **20% penalty** for **each** late day!
- Please carefully check the README of each assignment.

What to Do: Overview

You need to implement three major components of DDPMs:

- Forward Process: $q(\mathbf{x}_t | \mathbf{x}_0)$
- Reverse Process: $p_\theta(\mathbf{x}_{t-1} | \mathbf{x}_t)$
- Training Objective: $\|\epsilon - \epsilon_\theta(\mathbf{x}_t, t)\|^2$

Is it really that simple...?

What to Do: Overview

Yes! Even cutting-edge diffusion models are built this way.

By understanding this basic structure, you can begin exploring more advanced models.



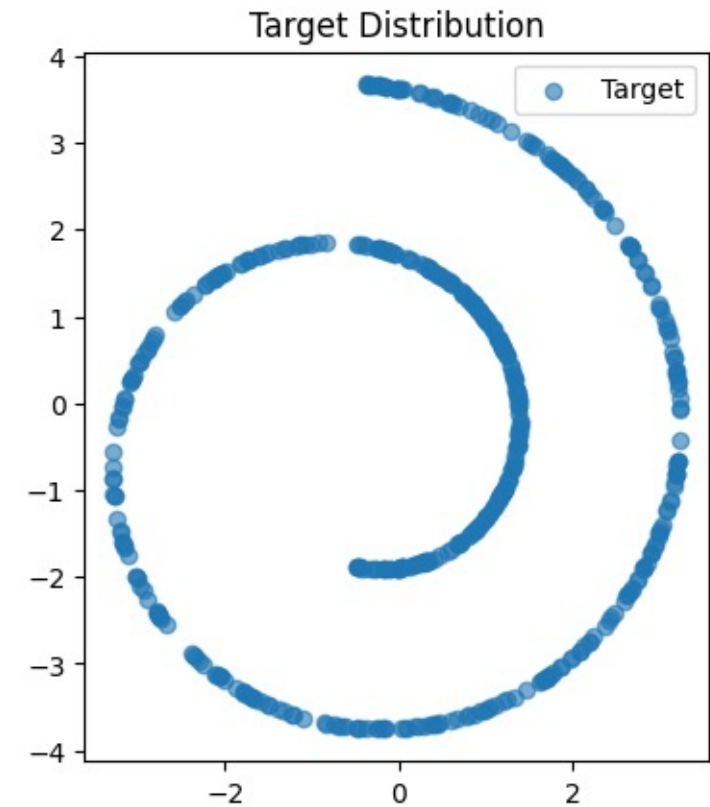

[Diffusers](#), HuggingFace



[Stable Diffusion 3](#), StabilityAI

What to Do: Task 1

Let's begin by modeling a simple distribution of 2D points (“Swiss Roll”).



What to Do: Task 1

Design a network that takes

- Noisy data x_t ;
- Current diffusion timestep t .

Hint: Use the `TimeLinear` class.

```
class SimpleNet(nn.Module):
    def __init__(
        self, dim_in: int, dim_out: int, dim_hids: List[int], num_timesteps: int
    ):
        super().__init__()
        """
        (TODO) Build a noise estimating network.

        Args:
            dim_in: dimension of input
            dim_out: dimension of output
            dim_hids: dimensions of hidden features
            num_timesteps: number of timesteps
        """

        ##### TODO #####
        # DO NOT change the code outside this part.

        #####

    def forward(self, x: torch.Tensor, t: torch.Tensor):
        """
        (TODO) Implement the forward pass. This should output
        the noise prediction of the noisy input x at timestep t.

        Args:
            x: the noisy data after t period diffusion
            t: the time that the forward diffusion has been running
        """

        ##### TODO #####
        # DO NOT change the code outside this part.

        #####

        return x
```

2d_plot_diffusion_todo/network.py

What to Do: Task 1

Implement functions

- `q_sample`;
- `p_sample`;
- `p_sample_loop`;
- `compute_loss`.

```
def q_sample(self, x0, t, noise=None):
    """
    sample x_t from q(x_t | x_0) of DDPM.

    Input:
        x0 (`torch.Tensor`): clean data to be mapped to timestep t in the forward process of DDPM.
        t (`torch.Tensor`): timestep
        noise (`torch.Tensor`, optional): random Gaussian noise. if None, randomly sample Gaussian noise in the function.
    Output:
        xt (`torch.Tensor`): noisy samples
    """
    if noise is None:
        noise = torch.randn_like(x0)

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Compute xt.
    alphas_prod_t = extract(self.var_scheduler.alphas_cumprod, t, x0)
    xt = x0

    #####

    return xt
```

2d_plot_diffusion_todo/ddpm.py


What to Do: Task 1

Check your implementation of `q_sample`!

Visualize $q(\mathbf{x}_t)$

```
fig, axes = plt.subplots(1, 10, figsize=(28, 3))
for i, t in enumerate(range(0, 500, 50)):
    x_t = ddpm.q_sample(target_ds[:num_vis_particles].to(device), (torch.ones(num_vis_particles) * t).to(device))
    x_t = x_t.cpu()
    axes[i].scatter(x_t[:,0], x_t[:,1], color='white', edgecolor='gray', s=5)
    axes[i].set_axis_off()
    axes[i].set_title('$q(\mathbf{x}_{'+str(t)+'})$')
```

[3] Python



2d_plot_diffusion_todo/ddpm_tutorial.ipynb

What to Do: Task 1

Implement functions

- `q_sample`;
- `p_sample`;
- `p_sample_loop`;
- `compute_loss`.

```
@torch.no_grad()
def p_sample(self, xt, t):
    """
    One step denoising function of DDPM:  $x_t \rightarrow x_{t-1}$ .

    Input:
        xt (`torch.Tensor`): samples at arbitrary timestep t.
        t (`torch.Tensor`): current timestep in a reverse process.
    Output:
        x_t_prev (`torch.Tensor`): one step denoised sample. (=  $x_{t-1}$ )

    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute x_t_prev.
    if isinstance(t, int):
        t = torch.tensor([t]).to(self.device)
    eps_factor = (1 - extract(self.var_scheduler.alphas, t, xt)) / (
        1 - extract(self.var_scheduler.alphas_cumprod, t, xt)
    ).sqrt()
    eps_theta = self.network(xt, t)

    x_t_prev = xt

    #####
    return x_t_prev
```

2d_plot_diffusion_todo/ddpm.py

What to Do: Task 1

Implement functions

- `q_sample`;
- `p_sample`;
- `p_sample_loop`;
- `compute_loss`.

```
@torch.no_grad()
def p_sample_loop(self, shape):
    """
    The loop of the reverse process of DDPM.

    Input:
        shape (`Tuple`): The shape of output. e.g., (num particles, 2)
    Output:
        x0_pred (`torch.Tensor`): The final denoised output through the DDPM reverse process.
    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # sample x0 based on Algorithm 2 of DDPM paper.
    x0_pred = torch.zeros(shape).to(self.device)

    #####
    return x0_pred
```

2d_plot_diffusion_todo/ddpm.py

What to Do: Task 1

Implement functions

- `q_sample;`
- `p_sample;`
- `p_sample_loop;`
- `compute_loss.`

```
def compute_loss(self, x0):
    """
    The simplified noise matching loss corresponding Equation 14 in DDPM paper.

    Input:
        x0 (`torch.Tensor`): clean data
    Output:
        loss: the computed loss to be backpropagated.
    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute noise matching loss.
    batch_size = x0.shape[0]
    t = (
        torch.randint(0, self.var_scheduler.num_train_timesteps, size=(batch_size,))
        .to(x0.device)
        .long()
    )



    loss = x0.mean()

    #####
    return loss
```

2d_plot_diffusion_todo/ddpm.py

What to Do: Task 1

Train your model and observe how the generated samples and the loss curve evolve over time.



What to Do: Task 2

We will now move on to a more interesting example: image generation.



Samples from our model trained using the [AFHQ dataset](#).

What to Do: Task 2

*Ugh...
Time to write more code...*



Samples from our model trained using the [AFHQ dataset](#).

What to Do: Task 2

Bring your codes from Task 1!

The code needs to be modified, but the changes should be kept minimal.

- `q_sample` → `add_noise`;
- `p_sample` → `step`;
- `compute_loss` → `get_loss`.

What to Do: Task 2

```
def q_sample(self, x0, t, noise=None):
    """
    sample x_t from q(x_t | x_0) of DDPM.

    Input:
        x0 (`torch.Tensor`): clean data to be mapped to timestep t in the forward process of DDPM.
        t (`torch.Tensor`): timestep
        noise (`torch.Tensor`, optional): random Gaussian noise. if None, randomly sample Gaussian noise in the function.
    Output:
        xt (`torch.Tensor`): noisy samples
    """
    if noise is None:
        noise = torch.randn_like(x0)

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Compute xt.
    alphas_prod_t = extract(self.var_scheduler.alphas_cumprod, t, x0)
    xt = x0

    #####

    return xt
```

2d_plot_diffusion_todo/ddpm.py

```
def add_noise(
    self,
    x_0: torch.Tensor,
    t: torch.IntTensor,
    eps: Optional[torch.Tensor] = None,
):
    """
    A forward pass of a Markov chain, i.e., q(x_t | x_0).

    Input:
        x_0 (`torch.Tensor [B,C,H,W]`): samples from a real data distribution q(x_0).
        t: (`torch.IntTensor [B]`)
        eps: (`torch.Tensor [B,C,H,W]`, optional): if None, randomly sample Gaussian noise in the function.
    Output:
        x_t: (`torch.Tensor [B,C,H,W]`): noisy samples at timestep t.
        eps: (`torch.Tensor [B,C,H,W]`): injected noise.
    """

    if eps is None:
        eps = torch.randn(x_0.shape, device='cuda')

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Assignment 1. Implement the DDPM forward step.
    x_t = None
    #####

    return x_t, eps
```

image_diffusion_todo/scheduler.py

What to Do: Task 2

```
@torch.no_grad()
def p_sample(self, xt, t):
    """
    One step denoising function of DDPM:  $x_t \rightarrow x_{t-1}$ .

    Input:
        xt (`torch.Tensor`): samples at arbitrary timestep t.
        t (`torch.Tensor`): current timestep in a reverse process.
    Output:
        x_t_prev (`torch.Tensor`): one step denoised sample. (=  $x_{t-1}$ )

    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute x_t_prev.
    if isinstance(t, int):
        t = torch.tensor([t]).to(self.device)
    eps_factor = (1 - extract(self.var_scheduler.alphas, t, xt)) / (
        1 - extract(self.var_scheduler.alphas_cumprod, t, xt)
    ).sqrt()
    eps_theta = self.network(xt, t)

    x_t_prev = xt

    #####
    return x_t_prev
```

2d_plot_diffusion_todo/ddpm.py

```
def step(self, x_t: torch.Tensor, t: int, eps_theta: torch.Tensor):
    """
    One step denoising function of DDPM:  $x_t \rightarrow x_{t-1}$ .

    Input:
        x_t (`torch.Tensor [B,C,H,W]`): samples at arbitrary timestep t.
        t (`int`): current timestep in a reverse process.
        eps_theta (`torch.Tensor [B,C,H,W]`): predicted noise from a learned model.
    Output:
        sample_prev (`torch.Tensor [B,C,H,W]`): one step denoised sample. (=  $x_{t-1}$ )

    """

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Assignment 1. Implement the DDPM reverse step.
    sample_prev = None
    #####

    return sample_prev
```

image_diffusion_todo/scheduler.py

What to Do: Task 2

```
def compute_loss(self, x0):
    """
    The simplified noise matching loss corresponding Equation 14 in DDPM paper.

    Input:
        x0 (`torch.Tensor`): clean data
    Output:
        loss: the computed loss to be backpropagated.
    """
    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute noise matching loss.
    batch_size = x0.shape[0]
    t = (
        torch.randint(0, self.var_scheduler.num_train_timesteps, size=(batch_size,))
        .to(x0.device)
        .long()
    )

    loss = x0.mean()

    #####
    return loss
```

2d_plot_diffusion_todo/ddpm.py

```
def get_loss(self, x0, class_label=None, noise=None):
    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute noise matching loss.
    B = x0.shape[0]
    timestep = self.var_scheduler.uniform_sample_t(B, self.device)
    loss = x0.mean()
    #####
    return loss
```


image_diffusion_todo/scheduler.py

What to Do: Task 2


After implementing the functions, start training the model by running

```
python train.py
```


The results will be saved under results directory.



Step 0



Step 1K



Step 25K

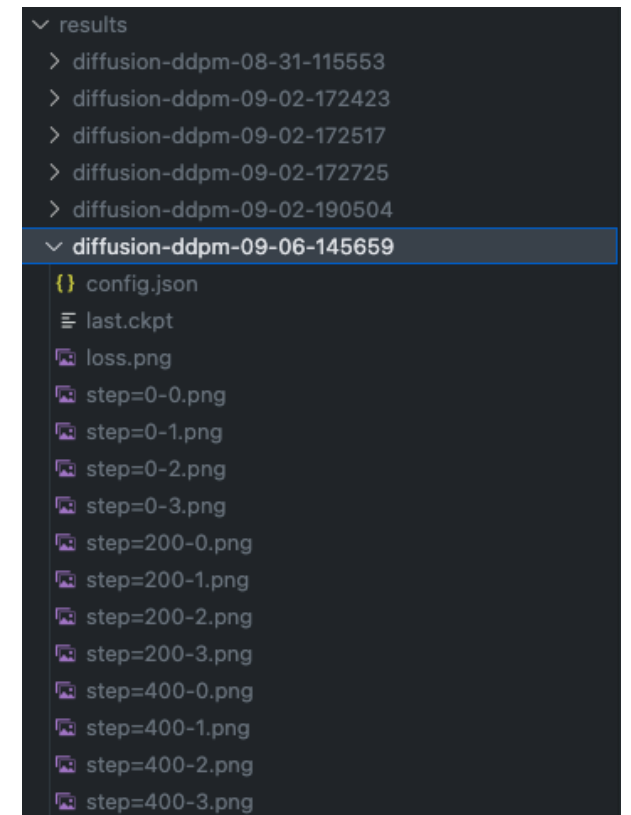


Step 50K



Step 100K


Images logged during training.



What to Do: Task 2

Generate the images using the trained model by running

```
python sampling.py \  
--ckpt_path {CKPT} --save_dir {SAVE}
```




What to Do: Task 2

Generate the images using the trained model by running

```
python sampling.py \  
--ckpt_path {CKPT} --save_dir {SAVE}
```



Samples generated using our model.




What to Do: Task 2

Prepare the data for evaluation by running

`python dataset.py` (Only once!)

This will create the `eval` directory under `data/afhq`.

Do NOT forget to run this. Otherwise, you will get incorrect FIDs!



FID: 229.33834138594412

FID: 10.844529455220403



FID scores across different test sets using the same generated samples.

What to Submit

Compile the following items into a PDF file: `{NAME}_{ID}.pdf`.

Task 1

- A screenshot of the loss curve;
- A screenshot of the Chamfer Distance;
- A visualization of samples generated using your DDPM.

Task 2

- A screenshot of the computed FID;
- At least 8 images generated using your DDPM.

What to Submit

Create a single ZIP file `{NAME}_{ID}.zip` including:

- The PDF file formatted following the guideline;
- Your code *without* checkpoints for DDPMs and the Inception Network

Your score will be deducted by 10% for each missing item.

Please check carefully!

Grading

You will receive up to 20 points from this assignment.

Task 1

- 10 points: Achieve CD lower than 20.
- 5 points: Achieve CD greater, or equal to 20 and less than 40.
- 0 point: Otherwise.

Grading

You will receive up to 20 points from this assignment.

Task 2

- 10 points: Achieve FID lower than 20.
- 5 points: Achieve FID greater, or equal to 20 and less than 40.
- 0 point: Otherwise.

Demo

Thank You