

CS492D: Diffusion Models and Their Applications

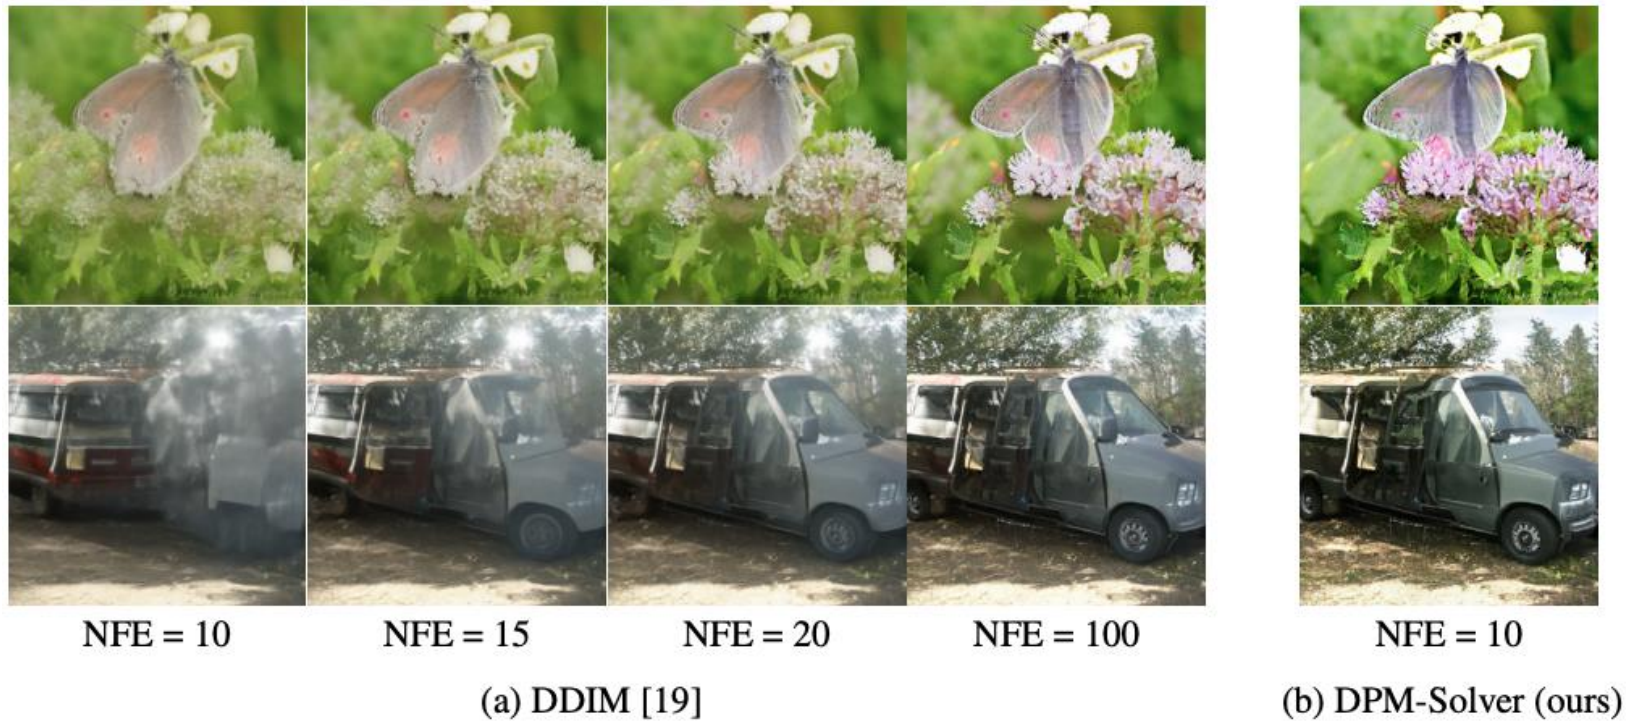
Assignment 6 Session

JUIL KOO

Fall 2024
KAIST

Introduction

In Assignment 6, we will implement **DPM-Solver**, a training-free approach to accelerate the reverse process using high-order ODE solvers.

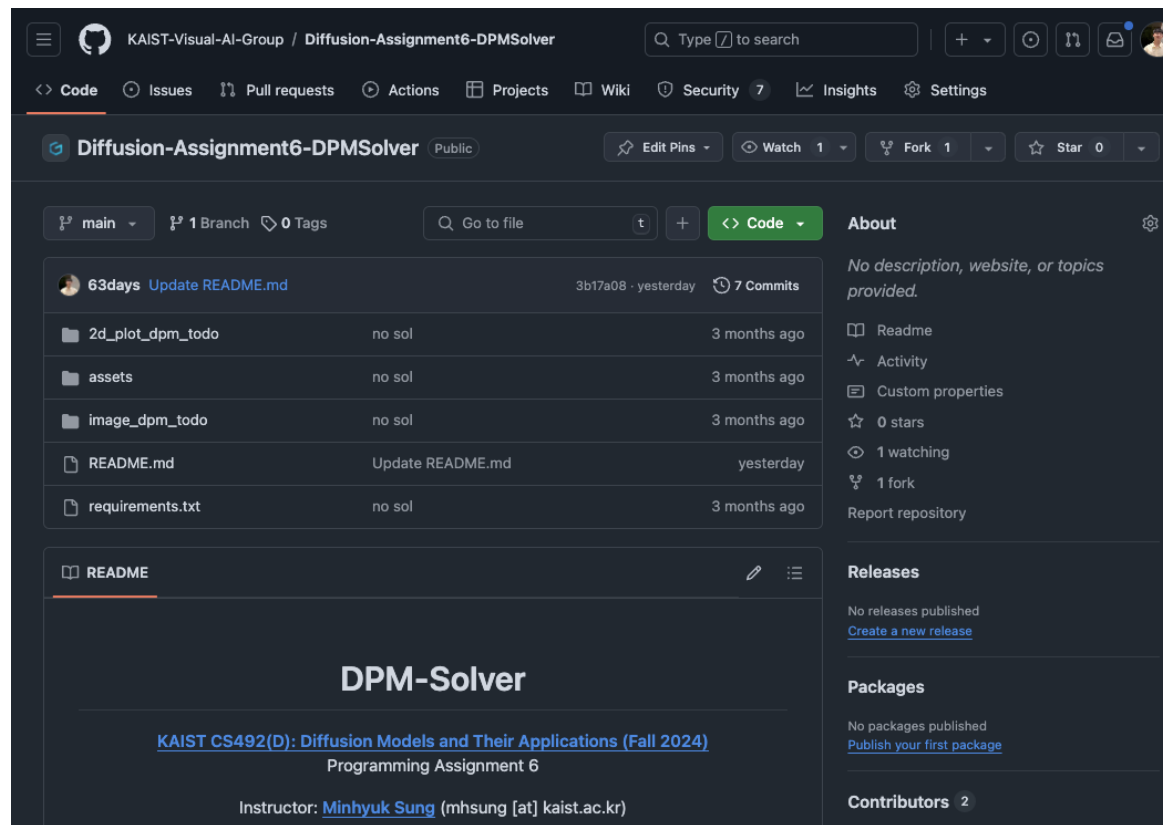


DPM-Solver, Lu *et al.*, NeurIPS 2022

Introduction

The skeleton code and instructions are available at:

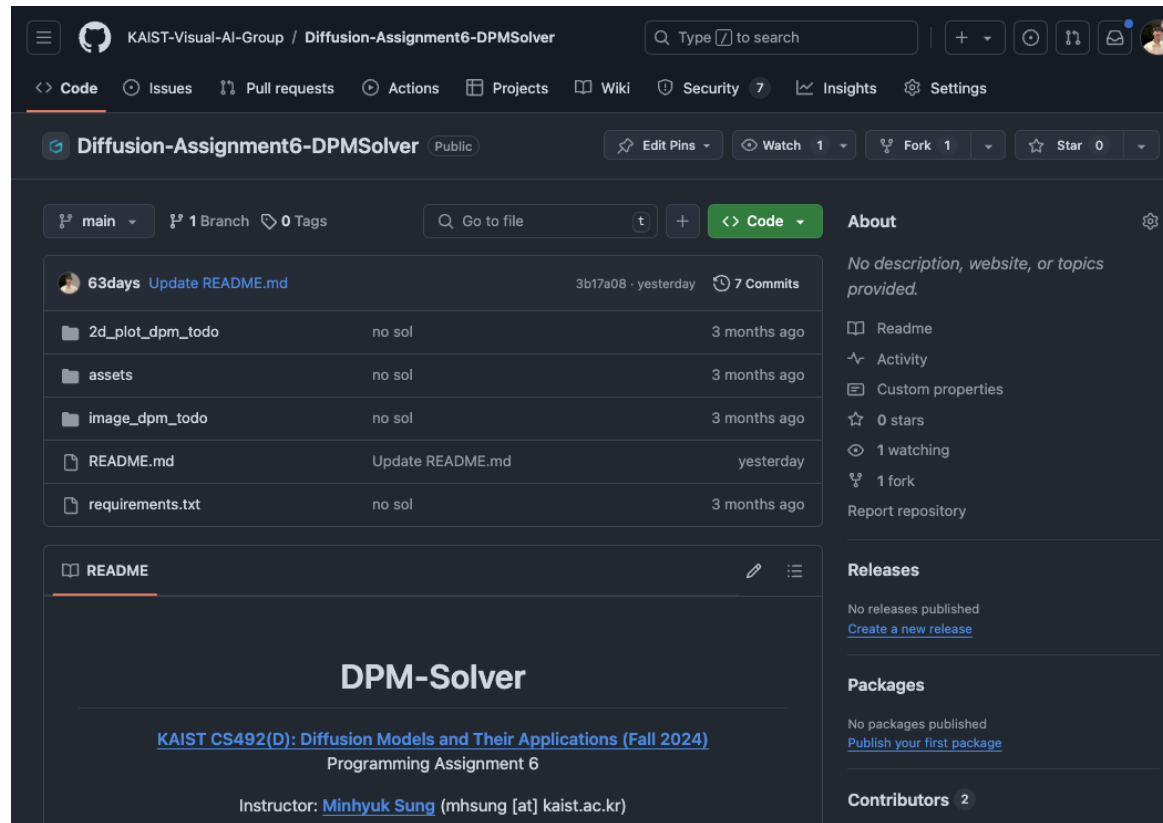
<https://github.com/KAIST-Visual-AI-Group/Diffusion-Assignment6-DPMSolver>



Introduction

Assignment 6 builds on your implementation of Assignment 2 (DDIM-CFG).

Refer to the `README.md` for details on which files and models need to be copied from Assignment 2.



Important Notes

- Assignment 6 is due on **Dec 6th (Friday)**.
- Late submission will incur **20% penalty** for **each** late day!
- Please carefully check the README of each assignment.
- Missing items in your submission will also incur penalties.

Connection Between DDPM and SDE

Forward Process: $q_{0t}(\mathbf{x}_t|\mathbf{x}_0) = \mathcal{N}(\mathbf{x}_t|\alpha(t)\mathbf{x}_0, \sigma^2(t)\mathbf{I})$

Corresponding SDE: $d\mathbf{x}_t = f(t)\mathbf{x}_t dt + g(t)d\mathbf{w}_t$,

where $f(t) = \frac{d \log \alpha(t)}{dt}$ and $g^2(t) = \frac{d\sigma^2(t)}{dt} - 2 \frac{d \log \alpha(t)}{dt} \sigma^2(t)$.

Continuous Forms of DDPM

Song *et al.* has shown the **forward process** has its equivalent **reverse** SDE and PF-ODE.

SDE of **Forward Process**: $d\mathbf{x}_t = f(t)\mathbf{x}_t dt + g(t)d\mathbf{w}_t$

SDE of **Reverse Process**: $d\mathbf{x}_t = [f(t)\mathbf{x}_t - g^2(t)\nabla_{\mathbf{x}} \log q_t(\mathbf{x}_t)]dt + g(t)d\bar{\mathbf{w}}_t$

PF-ODE of **Reverse Process**:
$$\begin{aligned} \frac{d\mathbf{x}_t}{dt} &= f(t)\mathbf{x}_t - \frac{1}{2}g^2(t)\nabla_{\mathbf{x}} \log q_t(\mathbf{x}_t) \\ &= f(t)\mathbf{x}_t + \frac{g^2(t)}{2\sigma_t} \boldsymbol{\epsilon}_{\theta}(\mathbf{x}_t, t) \end{aligned}$$

Song *et al.*, Score-Based Generative Modeling through Stochastic Differential Equations, ICLR 2021.

Exact Solutions of PF-ODEs

PF-ODE of **Reverse Process**: $\frac{d\mathbf{x}_t}{dt} = f(t)\mathbf{x}_t + \frac{g^2(t)}{2\sigma_t} \boldsymbol{\epsilon}_\theta(\mathbf{x}_t, t)$

The PF-ODE is a 1st-order semi-linear differential equation, whose solution can be computed. Starting from $s > t$, the solution at t is:

$$\mathbf{x}_t = \int_s^t \frac{d\mathbf{x}_\tau}{d\tau} d\tau = e^{\int_s^t f(\tau) d\tau} \mathbf{x}_s + \int_s^t \left(e^{\int_\tau^t f(r) dr} \frac{g^2(\tau)}{2\sigma_\tau} \boldsymbol{\epsilon}_\theta(\mathbf{x}_\tau, \tau) \right) d\tau.$$

Exact Solutions of PF-ODEs

$$\text{Exact solution: } \mathbf{x}_t = \underbrace{e^{\int_s^t f(\tau) d\tau} \mathbf{x}_s}_{\text{Linear term}} + \underbrace{\int_s^t \left(e^{\int_\tau^t f(r) dr} \frac{g^2(\tau)}{2\sigma_\tau} \boldsymbol{\epsilon}_\theta(\mathbf{x}_\tau, \tau) \right) d\tau}_{\text{Non-linear term}}$$

Exact Solutions of PF-ODEs

Exact solution: $\mathbf{x}_t = e^{\int_s^t f(\tau) d\tau} \mathbf{x}_s + \int_s^t \left(e^{\int_\tau^t f(r) dr} \frac{g^2(\tau)}{2\sigma_\tau} \boldsymbol{\epsilon}_\theta(\mathbf{x}_\tau, \tau) \right) d\tau$

The solution can be greatly simplified by introducing a new variable $\lambda_t := \log\left(\frac{\alpha_t}{\sigma_t}\right)$:

$$\mathbf{x}_t = \frac{\alpha_t}{\alpha_s} \mathbf{x}_s - \alpha_t \int_s^t \left(\frac{d\lambda_\tau}{d\tau} \right) \frac{\sigma_\tau}{\alpha_\tau} \boldsymbol{\epsilon}_\theta(\mathbf{x}_\tau, \tau) d\tau.$$

Rewrite the solution again by “*change-of-variable*” for $t \rightarrow \lambda_t$:

$\lambda_t = \lambda(t)$ is a **strictly decreasing** function of t .

$$\mathbf{x}_t = \frac{\alpha_t}{\alpha_s} \mathbf{x}_s - \alpha_t \int_{\lambda_s}^{\lambda_t} e^{-\lambda_\tau} \hat{\boldsymbol{\epsilon}}_\theta(\hat{\mathbf{x}}_{\lambda_\tau}, \lambda_\tau) d\lambda_\tau.$$

High-Order Solvers for PF-ODEs

Solution: $x_t = \frac{\alpha_t}{\alpha_s} x_s - \alpha_t \int_{\lambda_s}^{\lambda_t} e^{-\lambda} \hat{\epsilon}_\theta(\hat{\mathbf{x}}_\lambda, \lambda) d\lambda.$

Approximate the integral term by Taylor expansion of $\hat{\epsilon}_\theta(\hat{\mathbf{x}}_\lambda, \lambda)$:

$$\hat{\epsilon}_\theta(\hat{\mathbf{x}}_\lambda, \lambda) = \sum_{k=0}^{n-1} \frac{(\lambda - \lambda_s)^k}{k!} \hat{\epsilon}_\theta^{(k)}(\hat{\mathbf{x}}_{\lambda_s}, \lambda_s) + \mathcal{O}((\lambda - \lambda_s)^n).$$

n -th order derivative

$$\hat{\epsilon}_\theta^{(n)}(\hat{\mathbf{x}}_\lambda, \lambda) := \frac{d^n \hat{\epsilon}_\theta(\hat{\mathbf{x}}_\lambda, \lambda)}{d\lambda^n}$$

Feeding the Taylor expansion into the solution above results in:

$$x_t = \frac{\alpha_t}{\alpha_s} x_s - \alpha_t \sum_{k=0}^{n-1} \hat{\epsilon}_\theta^{(k)}(\hat{\mathbf{x}}_{\lambda_s}, \lambda_s) \int_{\lambda_s}^{\lambda_t} e^{-\lambda} \frac{(\lambda - \lambda_s)^k}{k!} d\lambda$$

This term only needs to be approximated.

DPM-Solver- n

$$x_t = \frac{\alpha_t}{\alpha_{t+1}} x_{t+1} - \alpha_t \sum_{k=0}^{n-1} \hat{\epsilon}_{\theta}^{(k)}(\hat{x}_{\lambda_{t+1}}, \lambda_{t+1}) \int_{\lambda_{t+1}}^{\lambda_t} e^{-\lambda} \frac{(\lambda - \lambda_{t+1})^k}{k!} d\lambda$$

(For clarity, we now change the notation s to $t + 1$.)

DPM-Solver- n is an approximation with the $(n - 1)$ -th order Taylor expansion.

DPM-Solver-1

$$x_t = \frac{\alpha_t}{\alpha_{t+1}} x_{t+1} - \alpha_t \sum_{k=0}^{n-1} \hat{\epsilon}_{\theta}^{(k)}(\hat{x}_{\lambda_{t+1}}, \lambda_{t+1}) \int_{\lambda_{t+1}}^{\lambda_t} e^{-\lambda} \frac{(\lambda - \lambda_{t+1})^k}{k!} d\lambda$$

$$\mathbf{x}_t = \frac{\alpha_t}{\alpha_{t+1}} \mathbf{x}_{t+1} - \sigma_t (e^{\lambda_t - \lambda_{t+1}} - 1) \epsilon_{\theta}(\mathbf{x}_{t+1}, t + 1)$$

$$= \frac{\alpha_t}{\alpha_{t+1}} \mathbf{x}_{t+1} - \alpha_t \left(\frac{\sigma_{t+1}}{\alpha_{t+1}} - \frac{\sigma_t}{\alpha_t} \right) \epsilon_{\theta}(\mathbf{x}_{t+1}, t + 1)$$

$$\because \lambda_t := \log\left(\frac{\alpha_t}{\sigma_t}\right)$$

DPM-Solver-1 is identical to DDIM.

DPM-Solver-2

1. **Split** the λ interval $(\lambda_t, \lambda_{t+1}]$ in half:

$$\tilde{\lambda} = \frac{\lambda_t + \lambda_{t+1}}{2}$$

2. Perform the **first-order approximation** for the first half:


$$\mathbf{u} \leftarrow \frac{\alpha_{t_\lambda(\tilde{\lambda})}}{\alpha_{t+1}} \mathbf{x}_{t+1} - \sigma_{t_\lambda(\tilde{\lambda})} \left(e^{\frac{h_t}{2}} - 1 \right) \hat{\boldsymbol{\epsilon}}_\theta(\mathbf{x}_{t+1}, t + 1)$$

3. Use the approximation to **update** $\hat{\boldsymbol{\epsilon}}_\theta$ and compute the final \mathbf{x}_t :

$$\mathbf{x}_t \approx \frac{\alpha_t}{\alpha_{t+1}} \mathbf{x}_{t+1} - \sigma_t (e^{h_t} - 1) \hat{\boldsymbol{\epsilon}}_\theta(\mathbf{u}, t_\lambda(\tilde{\lambda}))$$

Task 1

Implement DPM-Solver-1 for Swiss-Roll Data and 2D Images



Task 1

You need to fill in the parts marked with TODO in `scheduler.py`.

```
def first_order_step(self, x_s, s, t, eps_theta):
    """
    Implement Eq 4.1. in the DPM-Solver paper.
    Input:
        x_s (`torch.Tensor`): samples at timestep s.
        s (`torch.Tensor`): denoising starting timestep.
        t (`torch.Tensor`): denoising end timestep.
        eps_theta (`torch.Tensor`): noise prediction at s.
    Output:
        x_t (`torch.Tensor`): one step denoised sample.
    """
    assert torch.all(s > t), f"timestep s should be larger than timestep t"
    ##### TODO #####
    # DO NOT change the code outside this part.
    alpha_s = extract(self.dpm_alphas, s, x_s)
    x_t = x_s
    #####
    return x_t
```





Task 1

For 2D image generation, use a model that you trained with a CFG setup from Assignment 2.

Introduction

We also explore **Classifier-Free Guidance (CFG)**, a simple technique to enhance image quality in conditional generation.

"Pembroke Welsh corgi"



Weak Guidance Scale

Strong Guidance Scale

Diffusion Models Beat GANs on Image Synthesis, Dhariwal and Nichol, PMLR 2021

Task 2 (Optional)

Similarly, implement DPM-Solver-2 for Swiss-Roll Data and 2D Images.

You may want to access the noise prediction network's forward function inside a scheduler class.

```
class DPMSolverScheduler(BaseScheduler):
    def __init__(self, num_train_timesteps, beta_1=1e-4, beta_T=0.02, mode="linear"):
        assert mode == "linear", f"only linear scheduling is supported."
        super().__init__(num_train_timesteps, beta_1, beta_T, mode)
        self._convert_notations_ddpm_to_dpm()
        # To access the model forward in high-order scheduling.
        self.net_forward_fn = None
```

```
def build_dpm(config):
    network = SimpleNet(dim_in=2,
                        dim_out=2,
                        dim_hids=config["dim_hids"],
                        num_timesteps=config["num_diffusion_steps"]
                        )
    var_scheduler = DPMSolverScheduler(config["num_diffusion_steps"])
    var_scheduler.set_timesteps(config["num_inference_steps"])

    dpm = DiffusionModule(network, var_scheduler).to(config["device"])
    # For high-order sampling
    dpm.var_scheduler.net_forward_fn = dpm.network.forward
    return dpm

dpm = build_dpm(config)
```

Task 2 (Optional)

We also provide the `inverse_lambda()` function in `scheduler.py`, which corresponds to $t_\lambda(\lambda)$.

```
def inverse_lambda(self, lamb):
    """
    inverse function of lambda(t)
    """
    log_alpha_array = torch.log(self.dpm_alphas).reshape(1, -1)
    t_array = torch.linspace(0, 1, self.num_train_timesteps+1)[1:].reshape(1,-1).to(log_alpha_array)
    log_alpha = -0.5 * torch.logaddexp(torch.zeros((1,)).to(lamb.device), -2. * lamb)
    t = interpolate_fn(log_alpha.reshape((-1, 1)), torch.flip(log_alpha_array, [1]),
                      torch.flip(t_array, [1]))

    """
    Convert continuous t in [1 / N, 1] to discrete one in [0, 1000 * (N-1)/N]
    """
    t = ((t - 1 / self.num_train_timesteps) * 1000).long()
    return t.squeeze()
```

What to Submit

Include the following items into a PDF file: {NAME}_{SID}.pdf.

Task 1

- Sample Swiss-Roll data using DPM-Solver-1 and measure chamfer distance.
- Generate 500 images with the model trained in Assignment 2 using DPM-Solver-1, and measure FID.

Task 2 (Optional)

- Implement DPM-Solver-2 and repeat the same tasks in Task 1, but this time using DPM-Solver-2 instead of DPM-Solver-1.

What to Submit

Create a single ZIP file {NAME}_{SID}.zip including:

- The PDF file, formatted according to the guideline;
- Your implemented code.

Your score will be deducted by 10% for *each* missing item.

Please check carefully!

Grading

You will receive up to 10 points from this assignment.

For Task 1, you will receive:

- 10 points: Achieve CD lower than 40 on Swiss-Roll and achieve FID less than 30 on image generation with CFG.
- 5 points: Either achieve CD between 40 and 60 or achieve FID between 30 and 50.
- 0 points: Otherwise.

Grading: Compensation through Task 2

- **Task 2 is optional** and offers an additional 10 points.
- It follows the same criteria of Task 1, but must use DPM-Solver-2.
- These extra points can only compensate for **points lost in other assignments**, not in participation scores or projects.
- The total credits across all assignments will be 130 ($20 \times 6 + 10$).
(10 for Assignment 6 and 20 each for the others.)

Thank You