

CS492D: Diffusion Models and Their Applications


Assignment 7 Session

JUIL KOO


Fall 2024
KAIST

Introduction

In Assignment 7, we will implement **Flow Matching**, a novel generative model that offers straighter sampling trajectories than diffusion models.



Diffusion



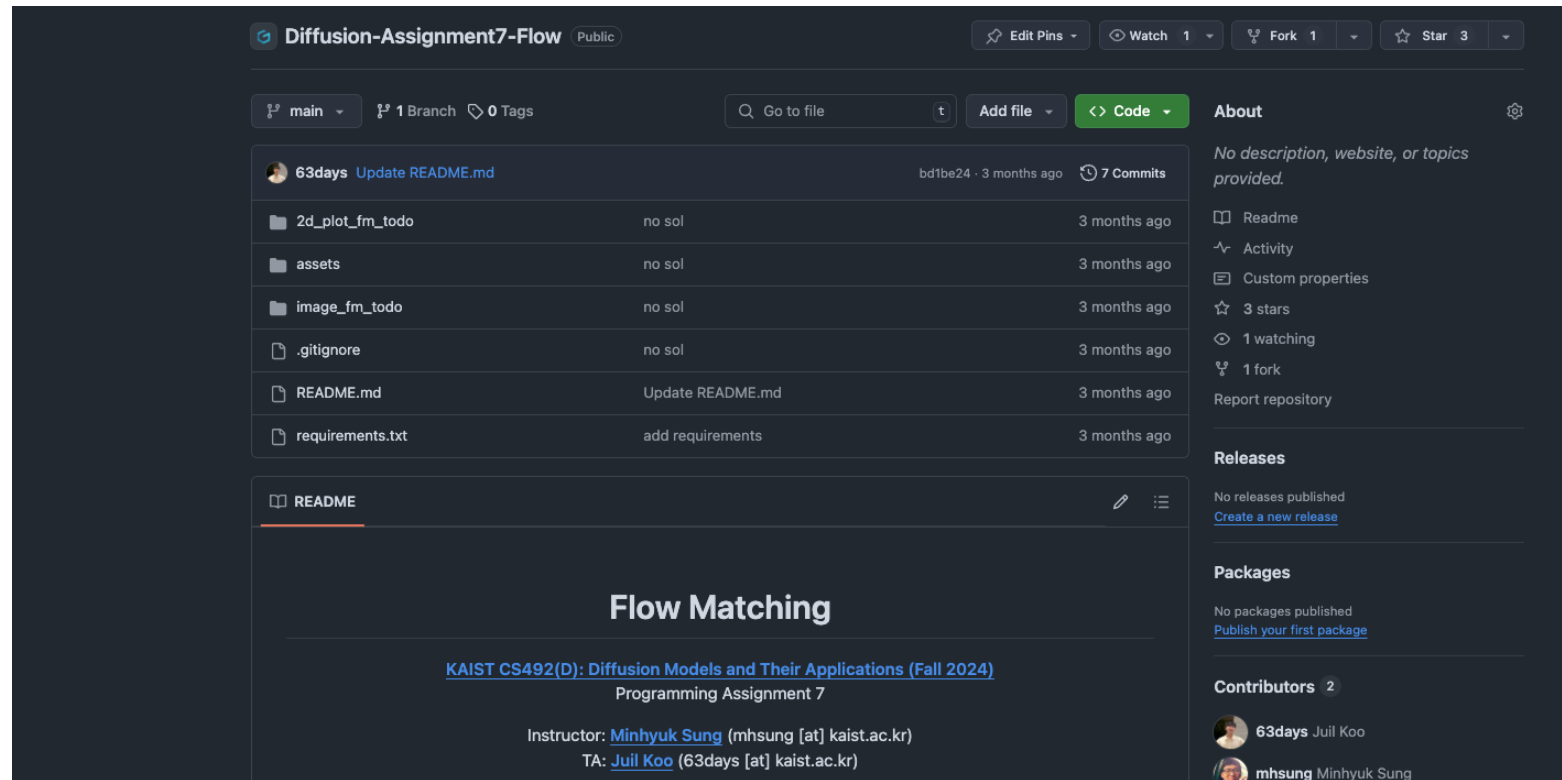
OT

Flow Matching for Generative Modeling, Lipman *et al.*, ICLR 2023.

Introduction


The skeleton code and instructions are available at:

<https://github.com/KAIST-Visual-AI-Group/Diffusion-Assignment7-Flow>



Introduction

Assignment 7 builds on your implementation of Assignment 2 (DDIM-CFG). Refer to the `README.md` for details on which files need to be copied.



Important Notes

- Assignment 7 is due on **Dec 13th (Friday)**.
- Late submission will incur **20% penalty** for **each** late day!
- Please carefully check the README of each assignment.
- Missing items in your submission will also incur penalties.

Notations

Diffusion Models [Generation: $T \rightarrow 0$]

- x_0 : A sample from the **data distribution**.
- x_T : A sample from the **reference (base) distribution**.

Flow Matching [Generation: $0 \rightarrow 1$]

- x_1 : A sample from the **data distribution**.
- x_0 : A sample from the **reference (base) distribution**.

Diffusion Models

Flow Matching

p_0



p_1



p_T

p_0

Flow and Vector Fields

- A flow $\phi_t(\mathbf{x})$ is a time dependent mapping function, which is similar to the forward pass of diffusion models: $\mathbf{x}_t = \phi_t(\mathbf{x})$.
- The derivative of $\phi_t(\mathbf{x})$ w.r.t t is called a vector field $u_t(\phi_t(\mathbf{x})) = \frac{d\phi_t(\mathbf{x})}{dt}$.

Probability Paths


A vector field $u_t(\phi_t(\mathbf{x}))$ is said to generate a probability path p_t if its flow $\phi_t(\mathbf{x})$ satisfies the equation below:

$$p_t(\mathbf{x}) = [\phi_t]_* p_0(\mathbf{x}) = p_0(\phi_t^{-1}(\mathbf{x})) \det \left[\frac{\partial \phi_t^{-1}}{\partial \mathbf{x}}(\mathbf{x}) \right]$$

Flow Matching

$$\mathcal{L}_{FM} = \mathbb{E}_{t, \mathbf{x}_t \sim p_t} [\|v_\theta(\mathbf{x}_t, t) - u_t(\mathbf{x}_t)\|^2]$$

However, we do not know how p_t and u_t look like. The vector field u_t that generates the probability path p_t is not unique.



<https://www.peterholderrieth.com/blog/2023/The-Fokker-Planck-Equation-and-Diffusion-Models/>

Conditional Flow Matching

Since the groundtruth vector field $u_t(\mathbf{x}_t)$ is unknown, we define p_t and u_t **per sample**, i.e., $p_{t|1}(\mathbf{x}_t|\mathbf{x}_1)$ and $u_{t|1}(\mathbf{x}_t|\mathbf{x}_1)$, where $\mathbf{x}_1 \sim p_1$.

$$\mathcal{L}_{FM} = \mathbb{E}_{t, \mathbf{x}_t \sim p_t} [\|v_\theta(\mathbf{x}_t, t) - u_t(\mathbf{x}_t)\|^2]$$



$$\mathcal{L}_{CFM} = \mathbb{E}_{t, \mathbf{x}_t \sim p_{t|1}} [\|v_\theta(\mathbf{x}_t, t) - u_{t|1}(\mathbf{x}_t|\mathbf{x}_1)\|^2]$$

Conditional $p_{t|1}$ and $u_{t|1}$

While the conditional probability paths $p_{t|1}$ and vector fields $u_{t|1}$ can be arbitrarily chosen, we adopt the simplest one, which is based on Gaussian distributions:

- **Conditional probability path $p_{t|1}$:**

$$p_{t|1}(\mathbf{x}_t|\mathbf{x}_1) = \mathcal{N}(\mu_t(\mathbf{x}_1), \sigma_t^2(\mathbf{x}_1)\mathbf{I})$$

- **Conditional flow map ϕ_t :**

$$\phi_t(\mathbf{x}_t|\mathbf{x}_1) = \mu_t(\mathbf{x}_1) + \sigma_t(\mathbf{x}_1)\mathbf{x}_t$$


- **Conditional vector field u_t :**

$$u_t(\mathbf{x}_t|\mathbf{x}_1) = \frac{\sigma'_t(\mathbf{x}_1)}{\sigma_t(\mathbf{x}_1)} (\mathbf{x}_t - \mu_t(\mathbf{x}_1)) + \mu'_t(\mathbf{x}_1)$$


Special Instance of FM: Straight Path

- Given a flow map $\phi_t(\mathbf{x}_t|\mathbf{x}_1) = \mu_t(\mathbf{x}_1) + \sigma_t(\mathbf{x}_1)\mathbf{x}_t$, we set $\mu_t(\mathbf{x}_1) = t\mathbf{x}_1$ and $\sigma_t(\mathbf{x}_1) = 1 - (1 - \sigma_{min})t$.
- This results in a simple interpolation between \mathbf{x}_1 and \mathbf{x}_t :

$$\phi_t(\mathbf{x}_t|\mathbf{x}_1) = (1 - (1 - \sigma_{min})t)\mathbf{x}_t + t\mathbf{x}_1.$$




Diffusion




OT

Special Instance of FM: Straight Path

$$\begin{aligned}\mathcal{L}_{CFM} &= \mathbb{E}_{t, \mathbf{x}_t \sim p_{t|1}} [\|v_{\theta}(\mathbf{x}_t, t) - u_{t|1}(\mathbf{x}_t | \mathbf{x}_1)\|^2] \\ &= \mathbb{E}_{t, \mathbf{x}_1, \mathbf{x}_0} [\|v_{\theta}(\mathbf{x}_t, t) - (\mathbf{x}_1 - (1 - \sigma_{min})\mathbf{x}_0)\|^2]\end{aligned}$$



Diffusion




OT

Task 0

Before starting the assignment, copy and paste your implementation of ``2d_plot_ddpm_todo/network.py`` and ``image_ddpm_todo/network.py`` from Assignment 2.

Task 1: Flow Matching with Swiss-Roll

As done in Assignment 1, 2, and 6, you will first implement FM and test it in a simple Swiss-Roll 2D distribution.



Task 1: Flow Matching with Swiss-Roll

Complete the functions `compute_psi_t()` and `step()` in `fm.py`.

```
def compute_psi_t(self, x1, t, x):
    """
    Compute the conditional flow  $\psi_t(x | x_1)$ .

    Note that time flows in the opposite direction compared to DDPM/DDIM.
    As  $t$  moves from 0 to 1, the probability paths shift from a prior distribution  $p_0(x)$ 
    to a more complex data distribution  $p_1(x)$ .

    Input:
        x1 (`torch.Tensor`): Data sample from the data distribution.
        t (`torch.Tensor`): Timestep in [0,1).
        x (`torch.Tensor`): The input to the conditional  $\psi_t(x)$ .

    Output:
        psi_t (`torch.Tensor`): The conditional flow at t.
    """
    t = expand_t(t, x1)

    ##### TODO #####
    # DO NOT change the code outside this part.
    # compute  $\psi_t(x)$ 

    psi_t = x1
    #####

    return psi_t
```

```
def step(self, xt, vt, dt):
    """
    The simplest ode solver as the first-order Euler method:
     $x_{\text{next}} = x_t + dt * v_t$ 
    """

    ##### TODO #####
    # DO NOT change the code outside this part.
    # implement each step of the first-order Euler method.
    x_next = xt
    #####

    return x_next
```





Task 2: Image Generation with FM

As done in Assignments 2 and 6, we will sample images using FM with a CFG setup.

Introduction

We also explore **Classifier-Free Guidance (CFG)**, a simple technique to enhance image quality in conditional generation.

"Pembroke Welsh corgi"



Weak Guidance Scale

Strong Guidance Scale

Diffusion Models Beat GANs on Image Synthesis, Dhariwal and Nichol, PMLR 2021

Task 2: Image Generation with FM

Finish implementing `sample()` and `get_loss()` functions to work with a CFG setup.

```
def get_loss(self, x1, class_label=None, x0=None):
    """
    The conditional flow matching objective, corresponding Eq. 23 in the FM paper.
    """
    batch_size = x1.shape[0]
    t = self.fm_scheduler.uniform_sample_t(batch_size).to(x1)
    if x0 is None:
        x0 = torch.randn_like(x1)

    ##### TODO #####
    # DO NOT change the code outside this part.
    # Implement the CFM objective.
    if class_label is not None:
        model_out = self.network(x1, t, class_label=class_label)
    else:
        model_out = self.network(x1, t)

    loss = x1.mean()
    #####

    return loss
```

What to Submit

Include the following items into a PDF file: {NAME}_{SID}.pdf.

Task 1

- Loss curvature screenshot
- Chamfer distance results of FM sampling with 50 inference steps.
- Visualization of FM sampling.

Task 2

- FID score result obtained with the CFG scale of 7.5.
- At least 8 images generated by Flow Matching.

What to Submit

Create a single ZIP file {NAME}_{SID}.zip including:

- The PDF file, formatted according to the guideline;
- Your implemented code.

Your score will be deducted by 10% for *each* missing item.

Please check carefully!

Grading

You will receive up to 20 points from this assignment.

For Task 1, you will receive:

- 10 points: Achieve CD lower than 40.
- 5 points: Achieve CD between 40 and 60.
- 0 points: Otherwise.

Grading

You will receive up to 20 points from this assignment.

For Task 2, you will receive:

- 10 points: Achieve FID between 30 with $\text{CFG}=7.5$.
- 5 points: Achieve FID between 30 and 50 with $\text{CFG}=7.5$.
- 0 points: Otherwise.

Thank You